

SECTION 1 / EXECUTIVE SUMMARY

CITED 40%

OF THE TIME.

LEADER: 80%.

Across 10 buyer questions tested against 4 answer engines, benchmarked against 3 competitors, this is where Claimzilla stands in AI-driven answers today. The gap is a visibility gap, fixable by the technical and content patterns listed below.

Methodology, engine sample, and the definition of every severity label are documented in Section 8 at the back of this report.

YOUR SHARE	LEADER SHARE	ENGINES	FINDINGS
40%	80%	4	33
+/- 8%, 95% CI	+/- 7%, 95% CI	ChatGPT · Claude · Gemini · Perplexity	7 high · 17 medium · 8 clean
Of 10 buyer questions tested	Across 3 competitors		

How to read “+/- N%, 95% CI”: CI stands for **Confidence Interval**. Each buyer question is sampled at N=3 independent trials per engine, so the cited-share band is computed from the trial counts (this run: n=120). The +/- half-width is the Wilson 95% confidence interval: if we re-ran the audit, the cited share would land inside this band 95 times out of 100. Smaller band = more sampling certainty.

THE TOP PATTERNS MOVING THE NUMBER

01	GEMINI (GOOGLE) NAMES THE TARGET DOMAIN FOR REPRESENTATIVE QUERIES	SEVERITY
	Cited in 1 / 10 buyer queries (0 errors). Gemini follows Google trust signals tightly; today claimzilla.example has weak SERP presence + zero AIO citations + a thin Knowledge Graph entry.	HIGH
02	LLMS.TXT DIRECTIVE FILE AT THE SITE ROOT	SEVERITY
	/llms.txt not present (HTTP 404). The file is the public spec for telling LLM crawlers what to index, what to summarize, and what to exclude; missing it means crawlers fall back to robots.txt heuristics and may sample from low-value pages.	HIGH
03	GOOGLE AI OVERVIEW SURFACES THE TARGET DOMAIN ON COMMERCIAL BUYER QUERIES	SEVERITY
	0 of 8 commercial buyer queries surfaced Claimzilla in the AI Overview; the AIO slot is held by Insurify (5 of 8), The Zebra (2 of 8), and NerdWallet (1 of 8) across the query set. AIO is the highest-conversion answer slot in Google search today.	HIGH

The full implementation roadmap (this month / this quarter / next quarter / beyond, each with its specific fix-hint and copy-paste remediation) is Section 5 of this PDF.

SECTION 2 / STATE OF AI CITATION FOR YOUR CATEGORY

4 ENGINES. 10 BUYER QUESTIONS PER ENGINE.

Every AI answer engine makes different trust decisions when a buyer asks about your category. We sent the 10 buyer questions you provided on intake to ChatGPT, Claude, Gemini and Perplexity and recorded who surfaced Claimzilla as a cited source.

CHATGPT	CLAUDE	GEMINI	PERPLEXITY
2 / 10 +/- 14%, 95% CI; n=30 cited ChatGPT queries	3 / 10 +/- 16%, 95% CI; n=30 cited Claude queries	1 / 10 +/- 11%, 95% CI; n=30 cited Gemini queries	4 / 10 +/- 17%, 95% CI; n=30 cited Perplexity queries

Across **10** buyer questions, the engines collectively cited Claimzilla **10 of 40** times. That asymmetry is the visibility gap Section 4 gives you the patterns to close.

Note: this section reads the same citation dataset as Section 1. Section 1 aggregates into one headline. Section 2 breaks it per engine so you can see where the gap is widest.

SECTION 3 / COMPETITIVE POSITION

WHO THE ENGINES ARE CITING INSTEAD OF YOU.

You named **3** competitors on intake. The citation gap below is from the same buyer-question set as Section 1 and Section 2. When a competitor is cited and Claimzilla is not, their site has a visibility advantage you can reverse. The patterns in Section 4 explain how.

01	INSURIFY.COM Benchmarked against Claimzilla on the 10 buyer questions in Section 1. The aggregate citation gap across all named competitors is shown in the Section 1 KPI strip.	ROLE
		PRIMARY BENCHMARK
02	THEZEBRA.COM Benchmarked against Claimzilla on the 10 buyer questions in Section 1. The aggregate citation gap across all named competitors is shown in the Section 1 KPI strip.	ROLE
		PEER BENCHMARK
03	NERDWALLET.COM/INSURANCE Benchmarked against Claimzilla on the 10 buyer questions in Section 1. The aggregate citation gap across all named competitors is shown in the Section 1 KPI strip.	ROLE
		PEER BENCHMARK

SECTION 4 / FINDINGS AND RECOMMENDATIONS

THE PATTERNS, BY THEME.

Your full technical and content audit, grouped under the twelve themes that determine whether AI answer engines cite you. The per-check detail (click-to-copy remediations, pre-filled test commands, validation links) lives in your Playbook (see Section 7).

32 FINDINGS · 1 CHECK NOT APPLICABLE TO YOUR ENGAGEMENT

CRAWL & INDEXABILITY 17 FINDINGS

01	LLMS.TXT DIRECTIVE FILE AT THE SITE ROOT /llms.txt not present (HTTP 404). The file is the public spec for telling LLM crawlers what to index, what to summarize, and what to exclude; missing it means crawlers fall back to robots.txt heuristics and may sample from low-value pages. Detail: see your Playbook → All Checks → llms.txt directive file at the site root.	SEVERITY HIGH
02	BREADCRUMBLIST SCHEMA visible breadcrumbs (nav[aria-label*=breadcrumb]) but no BreadcrumbList JSON-LD Detail: see your Playbook → All Checks → BreadcrumbList schema.	SEVERITY HIGH
03	CONTENT FRESHNESS SIGNALS 0/3 freshness signals fresh (visible=stale, jsonld=stale, header=none) Detail: see your Playbook → All Checks → Content freshness signals.	SEVERITY HIGH
04	ANSWER-FIRST PARAGRAPH IN THE FIRST 80 TO 110 WORDS 2/5 pages place a direct answer in the first paragraph; 3/5 lead with marketing copy or navigation prose Detail: see your Playbook → All Checks → Answer-first paragraph in the first 80 to 110 words.	SEVERITY HIGH
05	INLINE CITATIONS AND EXTERNAL-SOURCE FOOTNOTES 1/5 pages cite at least 2 external authoritative sources; 4/5 have zero outbound citations Detail: see your Playbook → All Checks → Inline citations and external-source footnotes.	SEVERITY HIGH

06	EFFECTIVE META TITLE title score 75% (52 chars); issues: title omits brand token 'claimzilla' <i>Detail: see your Playbook → All Checks → Effective meta title.</i>	SEVERITY MEDIUM
07	AUTHOR SCHEMA WITH LINKABLE BIOS 1/2 authors have a linkable bio (url or sameAs); names present but no bio link <i>Detail: see your Playbook → All Checks → Author schema with linkable bios.</i>	SEVERITY MEDIUM
08	SAMEAS ENTRIES ON ORGANIZATION SCHEMA 2 sameAs entries across 1 category; need 4+ entries; concentrated in one platform category <i>Detail: see your Playbook → All Checks → sameAs entries on Organization schema.</i>	SEVERITY MEDIUM
09	CONSISTENT BRAND NAME ACROSS SITE AND SCHEMA partial match: org JSON-LD 'Claimzilla Inc.' vs page 'Claimzilla' <i>Detail: see your Playbook → All Checks → Consistent brand name across site and schema.</i>	SEVERITY MEDIUM
10	E-E-A-T (EXPERIENCE, EXPERTISE, AUTHORITATIVENESS, TRUSTWORTHINESS) SIGNALS 3/5 E-E-A-T sub-signals present; missing: author bio, sources cited <i>Detail: see your Playbook → All Checks → E-E-A-T (Experience, Expertise, Authoritativeness, Trustworthiness) signals.</i>	SEVERITY MEDIUM
11	XML SITEMAP LASTMOD COVERAGE 67% lastmod coverage (42/63 URLs); 8/5 samples within 365d <i>Detail: see your Playbook → All Checks → XML sitemap lastmod coverage.</i>	SEVERITY MEDIUM
12	PRICING PAGE EXISTS AT THE SITE ROOT no /pricing path resolved; quote flow exists but no static price reference <i>Detail: see your Playbook → All Checks → Pricing page exists at the site root.</i>	SEVERITY MEDIUM
13	EFFECTIVE META DESCRIPTION description score 100% (148 chars)	SEVERITY CLEAN
14	ROBOTS.TXT DOES NOT BLOCK AI CRAWLERS 8/8 answer-engine bots allowed; training opt-out: Google-Extended	SEVERITY CLEAN

15

NOINDEX AI DIRECTIVES
11/11 AI bots can index this page

SEVERITY
CLEAN

16

PAGE LANGUAGE DECLARATION
lang='en-US'

SEVERITY
CLEAN

17

HOWTO SCHEMA ON STEP-BY-STEP GUIDES
no step-by-step content detected; HowTo schema not applicable for this page

SEVERITY
SKIP

STRUCTURED DATA

1 FINDING

01

JSON-LD: ORGANIZATION, WEBSITE, SERVICE, FAQPAGE SHAPES
Organization schema present and well-formed; WebSite + Service JSON-LD missing entirely; FAQPage schema NOT present despite 14 FAQ-shaped Q&A blocks on /faq.
Detail: see your Playbook → All Checks → JSON-LD: Organization, WebSite, Service, FAQPage shapes.

SEVERITY
MEDIUM

CONTENT STRUCTURE

1 FINDING

01

TITLE, DESCRIPTION, CANONICAL, ROBOTS, VIEWPORT, H1/H2, LANG
All 8 fundamentals present; title 58 chars, description 156 chars, canonical points to https://claimzilla.example/, single H1, robots index,follow, viewport configured, lang en-US.

SEVERITY
CLEAN

E-E-A-T & AUTHORITY

7 FINDINGS

01

GEMINI (GOOGLE) NAMES THE TARGET DOMAIN FOR REPRESENTATIVE QUERIES
Cited in 1 / 10 buyer queries (0 errors). Gemini follows Google trust signals tightly; today claimzilla.example has weak SERP presence + zero AIO citations + a thin Knowledge Graph entry.
Detail: see your Playbook → All Checks → Gemini (Google) names the target domain for representative queries.

SEVERITY
HIGH

02

OPENAI (CHATGPT) NAMES THE TARGET DOMAIN FOR REPRESENTATIVE QUERIES

SEVERITY

MEDIUM

Cited in 2 / 10 buyer queries for domain claimzilla.example (0 query errors). Strong on brand-name queries; weak on category queries where Insurify and The Zebra dominate the answer.

Detail: see your Playbook → All Checks → OpenAI (ChatGPT) names the target domain for representative queries.

03

ANTHROPIC (CLAUDE) NAMES THE TARGET DOMAIN FOR REPRESENTATIVE QUERIES

SEVERITY

MEDIUM

Cited in 3 / 10 buyer queries (0 errors). Best engine for Claimzilla today: Claude weights Wikidata + Wikipedia presence highly and claimzilla.example has both, lifting brand-recognition signals.

Detail: see your Playbook → All Checks → Anthropic (Claude) names the target domain for representative queries.

04

PERPLEXITY (SONAR) NAMES THE TARGET DOMAIN FOR REPRESENTATIVE QUERIES

SEVERITY

MEDIUM

Cited in 4 / 10 buyer queries (0 errors). Best citation rate of the four engines; Perplexity surfaces well-structured comparison content and Claimzilla's quote-flow pages match that pattern.

Detail: see your Playbook → All Checks → Perplexity (Sonar) names the target domain for representative queries.

05

WIKIPEDIA ARTICLE + EXTERNAL CITATIONS

SEVERITY

MEDIUM

Article exists (240 words, created 2023) but is below the 600-word threshold for stable LLM grounding. 2 external citations from third-party publications; competitors range from 8 to 21 external cites.

Detail: see your Playbook → All Checks → Wikipedia article + external citations.

06

GOOGLE KNOWLEDGE GRAPH ENTITY ENTRY

SEVERITY

MEDIUM

Knowledge Graph entry exists at /g/11h_xb27q3 (machine ID assigned) with name + URL + image + brief description. Missing: official social profiles, sameAs links, alternateName, founder, industry classification.

Detail: see your Playbook → All Checks → Google Knowledge Graph entity entry.

07

WIKIDATA Q-ID + TYPED CLAIMS (INSTANCE OF, INDUSTRY, FOUNDED)

SEVERITY

CLEAN

Q-ID Q12399847 exists with 11 typed claims including instance-of (insurance broker), industry (insurance), founded (2019), headquarters (Austin, Texas), and 5 social-media identifiers. Strong feed for Anthropic's entity-graph citation logic.

RENDERING & VISIBILITY

3 FINDINGS

01

GOOGLE AI OVERVIEW SURFACES THE TARGET DOMAIN ON COMMERCIAL BUYER QUERIES

0 of 8 commercial buyer queries surfaced Claimzilla in the AI Overview; the AIO slot is held by Insurify (5 of 8), The Zebra (2 of 8), and NerdWallet (1 of 8) across the query set. AIO is the highest-conversion answer slot in Google search today.

Detail: see your Playbook → All Checks → Google AI Overview surfaces the target domain on commercial buyer queries.

SEVERITY

HIGH

02

BING COPILOT ANSWER + BING ORGANIC RANK FOR BUYER QUERIES

Bing organic rank 4-7 across the 8 buyer queries (no brand-mention in Copilot answers); competitors hold ranks 1-3 consistently. Bing is the second-largest source of LLM grounding traffic after Google's index.

Detail: see your Playbook → All Checks → Bing Copilot answer + Bing organic rank for buyer queries.

SEVERITY

MEDIUM

03

GOOGLE ORGANIC RANK FOR BRAND-ONLY AND BRAND+MODIFIER QUERIES

Ranks 1 for `claimzilla` (brand-only); top 3 for `claimzilla insurance` and `claimzilla reviews`. Knowledge Panel renders for the brand-only query.

SEVERITY

CLEAN

MEDIA

2 FINDINGS

01

EMBEDDED YOUTUBE VIDEOS HAVE TRANSCRIPTS

5 YouTube embeds detected across the audited URLs; 2 carry <track> caption files; the other 3 (homepage hero, /how-it-works explainer, /reviews testimonial reel) ship without captions or transcript text.

Detail: see your Playbook → All Checks → Embedded YouTube videos have transcripts.

SEVERITY

MEDIUM

02

BRAND MENTIONS IN THIRD-PARTY YOUTUBE CONTENT (TITLES, DESCRIPTIONS, TRANSCRIPTS)

Brand mentioned in 12 third-party videos (combined ~140K views); competitors range 35-180 third-party videos. Strongest channels for Claimzilla: 'The Personal Finance Lab', 'Two Cents PBS', 'TechCrunch'.

Detail: see your Playbook → All Checks → Brand mentions in third-party YouTube content (titles, descriptions, transcripts).

SEVERITY

MEDIUM

SECURITY & TRUST

1 FINDING

01

GOOGLE BUSINESS PROFILE COMPLETENESS

Profile claimed at maps.google.com/place/Claimzilla but only 8 of 14 sections complete: hours present, phone present, website present, but services list missing, attributes empty, posts feed empty, Q&A unanswered (3 buyer questions sitting >30 days), photos limited (only logo + one office shot).

Detail: see your Playbook → All Checks → Google Business Profile completeness.

SEVERITY

MEDIUM

META & SOCIAL

1 FINDING

01

OPEN GRAPH + TWITTER CARD META COMPLETENESS

All 9 OG tags + 4 Twitter Card tags present and validate against OpenGraph.xyz and LinkedIn Post Inspector; share previews render correctly across Facebook / LinkedIn / Slack / iMessage.

SEVERITY

CLEAN

SECTION 6 / SCHEMA PACK SUMMARY

YOUR SCHEMA COVERAGE, AT A GLANCE.

After analyzing 25 pages of your site for schema opportunities, we found **1 in-use** and **3** recognized types still to add. The page-by-page breakdown is below. The ready-to-paste JSON-LD lives in one self-contained file per audited page inside your Playbook.

25%

YOUR SCHEMA COVERAGE

PARTIAL

STATE

Coverage 25% = **1 of 4** priority schema types present. 1 found, 3 missing.

PRESENT 1 / 4

Organization

MISSING 3 / 4

- WebSite
- Service
- FAQPage

The ready-to-paste schema code for every priority page lives in your Playbook. Open the Schema Pack section there for per-page drill-ins with copy buttons, validation links, and a "where to start" ordering.

Scope note: the PDF shows the coverage summary; the Playbook shows how to execute fixes.

SECTION 7 / IMPLEMENTATION HANDOFF

WHAT'S IN YOUR DELIVERY.

This PDF is the readout: every check, every finding, every recommendation. The schema-pack .zip attached to your delivery email holds ready-to-paste JSON-LD blocks for the structured data work, populated with Claimzilla's values. Where this PDF surfaces a Terminal command directly (Section 4 crawler-visibility verification, when applicable), copy it as-is to confirm the AI-crawler view from your machine.

Your interactive Playbook companion (click-to-copy remediations, validation links, full schema JSON per audited page) ships with this engagement; the Playbook link is in your delivery email.

30 days of async follow-up is included (email hello@web-cited.com).

HOW THIS AUDIT WAS RUN.

Every number in this report is reproducible from the data window stamped on the cover. This appendix documents the engines we sampled, the questions we asked, and how citation share was scored, so your team can validate, rerun, or challenge any result.

THE FOUR ENGINES WE TESTED

CHATGPT OpenAI Largest consumer and developer surface. Cites web sources when browsing is engaged or when responses draw from a retrieval index.	CLAUDE Anthropic Preferred by enterprise buyers. Citations surface through native web search tools and Claude.ai research sessions.	GEMINI Google Answer traffic overlaps with Google Search and AI Overviews. Citations follow Google's trust signals: entities, reviews, freshness.	PERPLEXITY Perplexity Answer engine with citations as a first-class product feature. Fastest signal for whether your pages are crawlable and quotable.
---	--	--	---

HOW CITATION SHARE IS COMPUTED

For each buyer question on the intake, we send the same prompt to every engine and capture the cited source list. A question is counted as **cited** when at least one of the engine responses links to Claimzilla's domain. The headline number on the cover is the percentage of questions where the engine collective cited you at least once, computed as $(all_cite + disagree) / total_queries$. That definition rewards any citation presence and penalizes only questions where no engine cited you at all.

THE SEVERITY LADDER

- 01

High: a failing signal that directly blocks citation, indexability, or trust. Fix this month.
- 02

Medium: a partial or inconsistent signal that reduces citation probability. Fix this quarter.
- 03

Clean: the signal is present and well-formed. No action required. Re-audit in six months to catch drift.
- 04

Info: observation that informs strategy but is not scored. Treat as context, not a task.

05

Skip: the check could not run (no data, unreachable page, or a dependency was absent). Listed for transparency.

SCOPE AND LIMITS

The audit covers the twelve themes listed in Section 4. Outside of scope for this report: paid search placements, sponsored content, vertical search engines (App Store, LinkedIn, YouTube comments), and any engine that does not return citation metadata. These are tracked separately in the Playbook where applicable.

COMPETITORS BENCHMARKED

The citation gap reported across this audit was measured against the 3 competitors named on intake: insurify.com, thezebra.com, nerdwallet.com/insurance. Each appears in the four LLM (Large Language Model) engines' answer text if and only if the engine surfaced that domain when answering the buyer questions in Section 1. To change the benchmark set on a re-audit, email hello@web-cited.com with an updated competitor list.

Data window: 2026-04-27 to 2026-04-27. Report generated 2026-05-06. Engagement #claimzilla-sample.

WEB CITED

Data collected 2026-04-27 ·
Reported 2026-05-06

#claimzilla-sample · SX0 Audit